

## EECS 388 – Winter 2015

### Lecture 21: Access Control and Isolation

#### Access Control

General setting: need to control access to resources (malicious programs, vulnerable programs)

subject: user, process, entity accessing resources

object: file, device, etc. being accessed

#### Access Control Matrix

	file1	file2
Alice	read	write
Bob	-	read
Charlie	-	-

#### Principle: "Complete Mediation"

control all resource requests through mediator that checks access control policy

#### Specifying policies:

##### ACL (access control lists):

- list associated w/ object
- user checked against list
- need to authenticate user

##### Capabilities:

- unforgeable ticket
- maybe can be passed from process to process
- ticket checker does not need to know user's id

delegation?

- process run under current user
- process can pass capability at runtime

revocation?

- remove user/group from list
- possible if bookkeeping

#### Unix file permissions:

- Processes associated w/ user IDs (can change)
- Files have permission bits, simple ACL:

setid	owner	group	other	
-	rwX	rwX	rwX	} vector of 4 octal value (e.g., chmod file 0644)
- Owner, root can change permissions
- if user == owner => owner permissions
- if user in group => group permissions
- else other permission

#### Capabilities

Instead of root vs. non-root, allow unprivileged processes to do whitelist of things

CAP\_NET\_ADMIN, CAP\_MKNOD, CAP\_SYS\_NICE, CAP\_SYS\_RAWIO

## Isolation

Traditional thinking: How to run bad/untrustworthy programs safely?

New thinking: Should be skeptical of all programs, isolate all you can!

Confinement: Keep app from harming rest of system

## Implementations:

air gap

[diagram]

virtual machines

[diagram]

syscall interposition

[diagram]

## Traditional OS mechanisms:

### Process isolation

Separate memory address spaces, explicit IPC

### chroot

- used, e.g., for guest access to ftp sites

chroot /tmp/guest -- root dir / is mapped to /tmp/guest

su guest -- Process ID becomes guest

- apps can't access files outside of jail

- utilities (ls, ps) must live in jail

- network access? (not restricted)

raw disk devices?

signals to other processes?

### FreeBSD jails

- hardened chroot

- can only bind to authorized sockets

- only communicate w/ processes inside jail

Not all programs can run in a jail (e.g., web browser cannot)

### Problems:

- coarse policies

- all-or-nothing filesystem access

- malicious apps can access netowkr, crash host OS

## System call interposition

- to damage host system, app makes system calls
  - filesystem: unlink, open, write
  - network: socket, bind, connect, send
- so, monitor syscalls, block unauthorized calls

### Linux: ptrace

complications:

- forks? fork monitor
- monitor crash? kill app
- must maintain state

ptrace isn't granular -- all system calls or none

security issues:

process 1: open ("me")

monitor: OK

check is not atomic

process 2: link me -> /etc/passwd

OS executes open("me")

"Time of check vs. time of use" vulnerability (TOCTOU)

Modern variants: (avoid TOCTOU correctly, but how to specify the right policies?)

fine-grained access control + system call interposition

policy defined in external file

e.g., App Armor: profiles app during learning phase, restricts privileges not in profile

e.g., SE Linux: similar, NSA developed

## Virtual Machines

apps	apps
guest OS	guest OS
VMM	
host OS	
hardware	

security assumptions:

- malware can infect guest OS, guest apps, but can't escape VM to host OS, other VMs
- problems: covert channels, etc.
  - malware might detect it's running inside VM, behave differently
  - have to monitor, manage, secure many OS instances

## Networks:

Physical isolation, VLANs, Firewalls

## Sandboxes:

Apply mechanisms like above to individual app or code

Separate tabs, SOP, etc. (browsers becoming more like OS)

Javascript in browser: API sandbox (no calls to do "dangerous" things)

Native client: Run untrusted x86 code in browser (Chrome)

Android, iOS application sandboxes, code signing